Árboles de decisión y ensambladores

Classification trees

	categorical	categorical	continuous	class
Tid	**Refund**	**Marital Status**	**Taxable Income**	**Cheat**
1	Yes	Single	125K	No
2	No	Married	100K	No
3	No	Single	70K	No
4	Yes	Married	120K	No
5	No	Divorced	95K	Yes
6	No	Married	60K	No
7	Yes	Divorced	220K	No
8	No	Single	85K	Yes
9	No	Married	75K	No
10	No	Single	90K	Yes

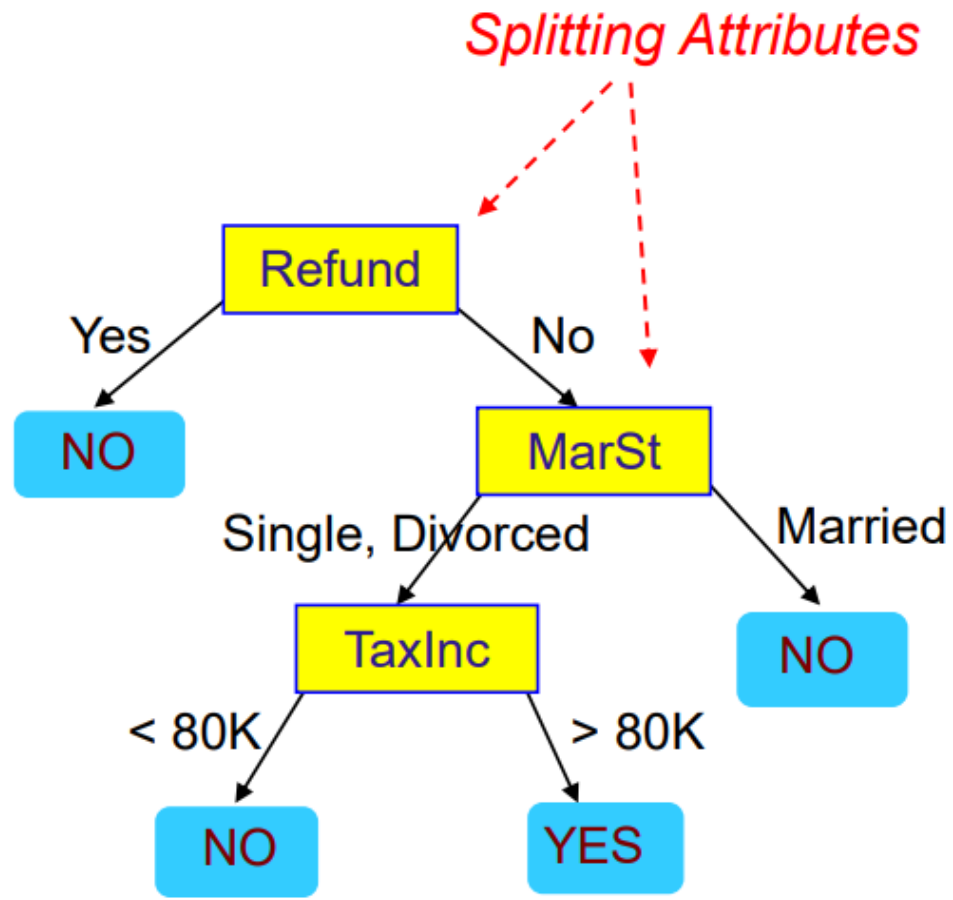

How to use

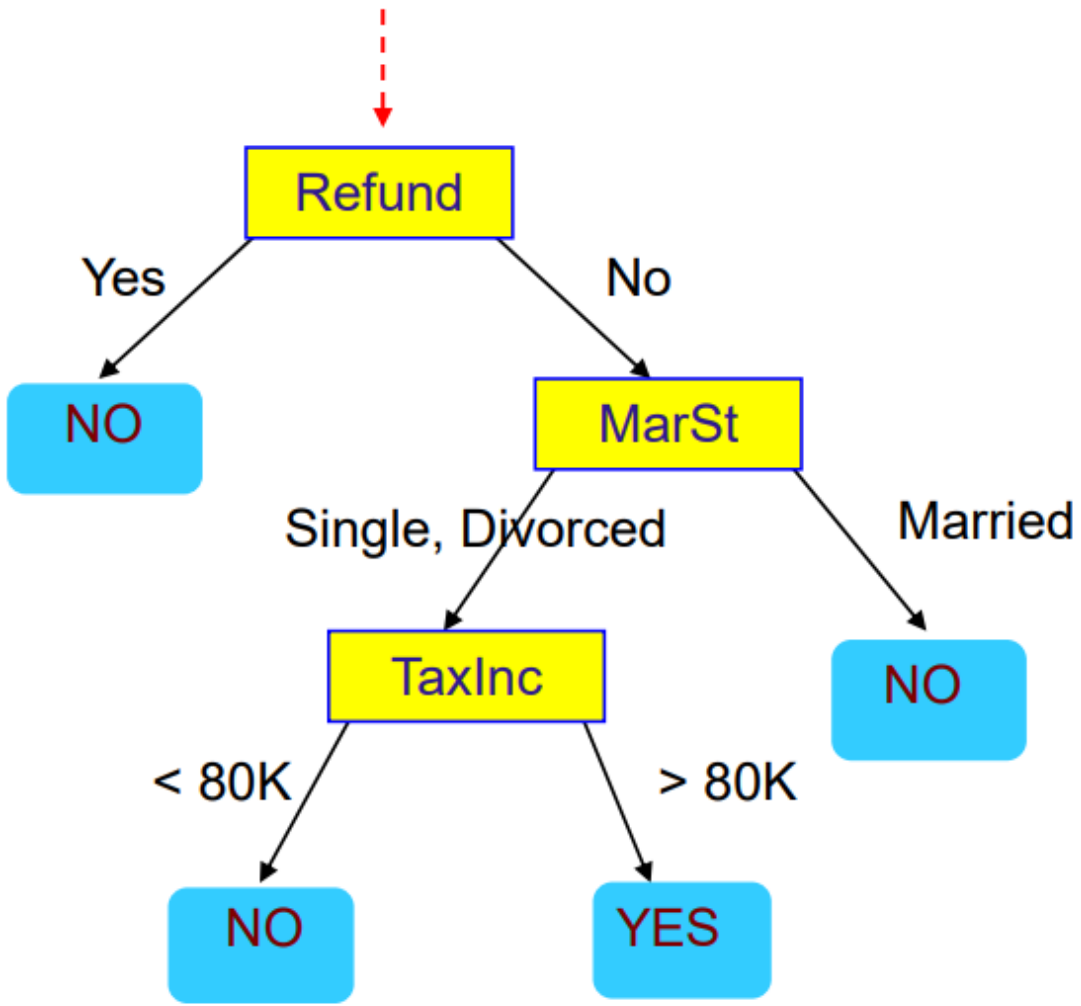

Refund	Marital Status	Taxable Income	Cheat
No	Married	80K	?

Aplicando el modelo a los datos de evaluación

Refund	Marital Status	Taxable Income	Cheat
No	Married	80K	?

Test Data

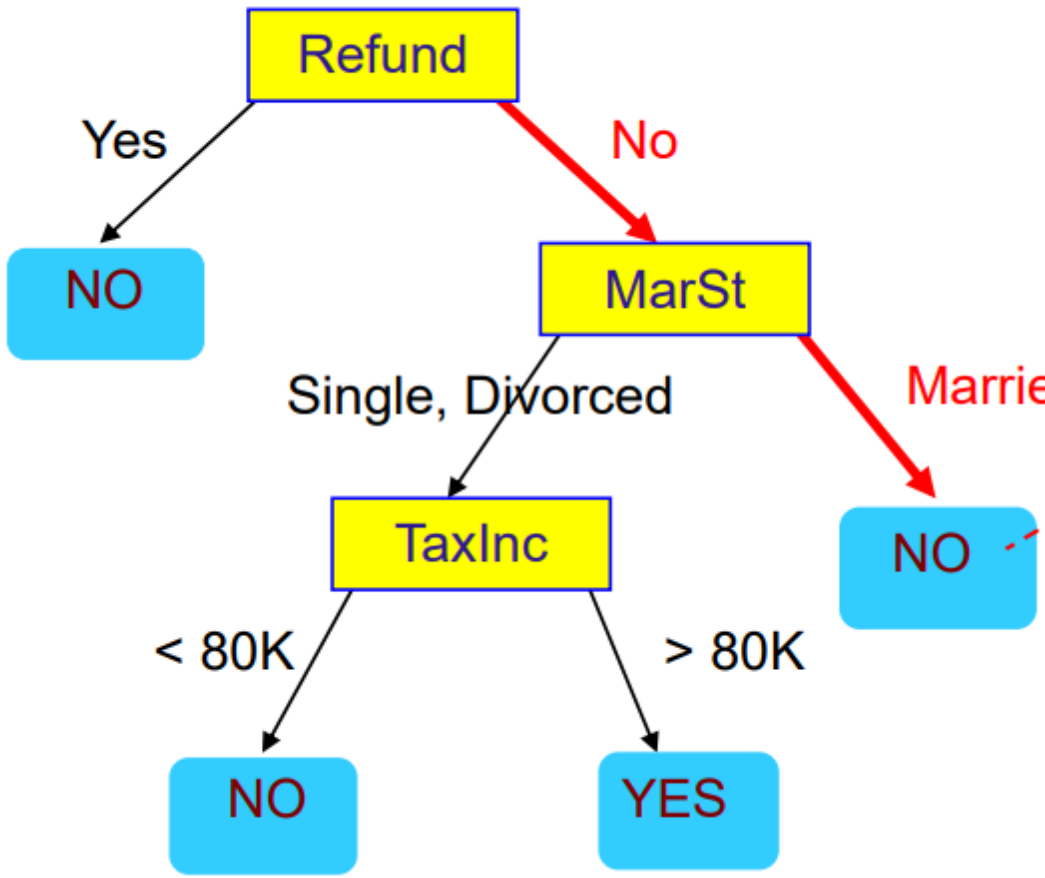

Assign Cheat to “No”

Impurity Criterion

Gini Index

$$I_G = 1 - \sum_{j=1}^{c} p_j^2$$

p_j: proportion of the samples that belongs to class c for a particular node

Entropy

$$I_H = -\sum_{j=1}^{c} p_j log_2(p_j)$$

p_j: proportion of the samples that belongs to class c for a particular node.

*This is the the definition of entropy for all non-empty classes ($p \neq 0$). The entropy is 0 if all samples at a node belong to the same class.

$$E(S) = \sum_{i=1}^{c} - p_i \log_2 p_i$$

Play Golf	
Yes	No
9	5

Entropy(PlayGolf) = Entropy (5,9)

= Entropy (0.36, 0.64)

= - (0.36 $\log_2$ 0.36) - (0.64 $\log_2$ 0.64)

= 0.94

Bagging y boosting

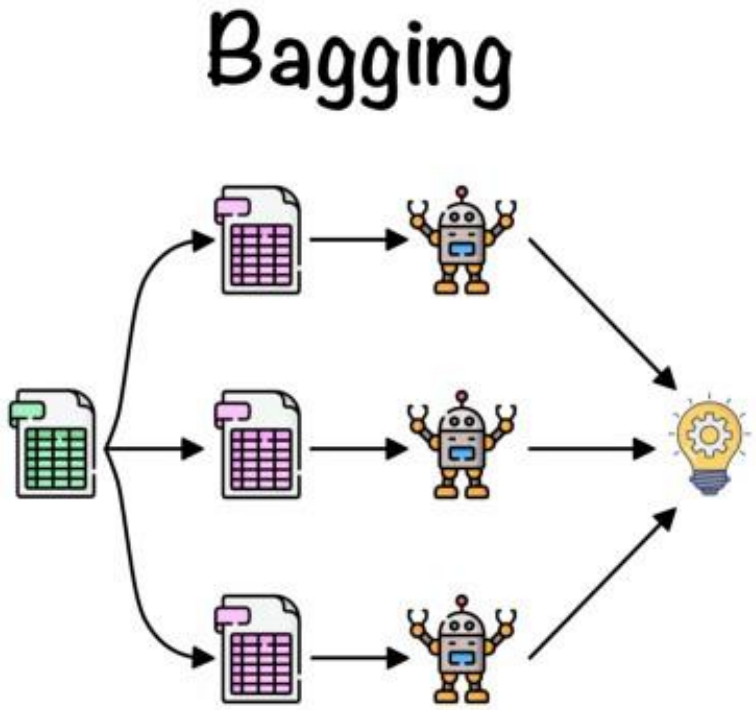

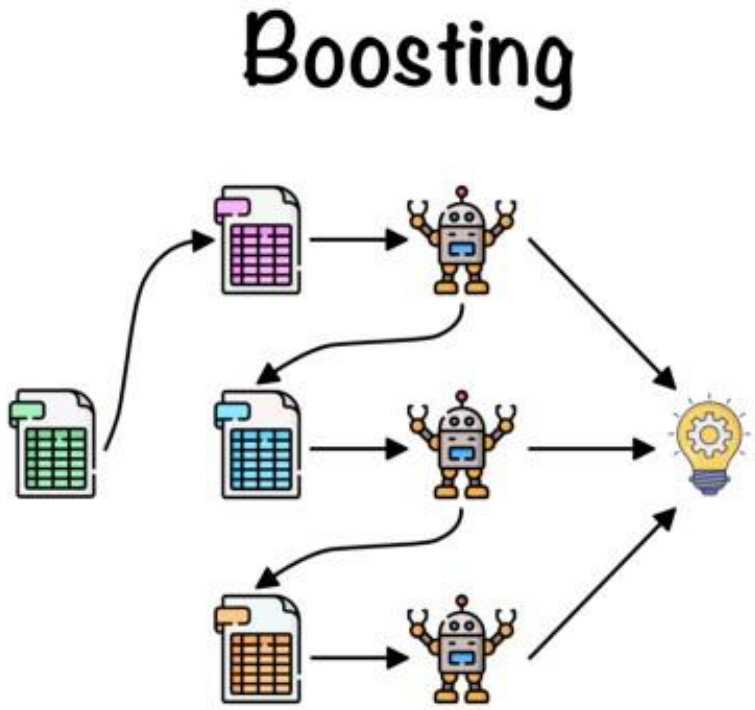

Bagging

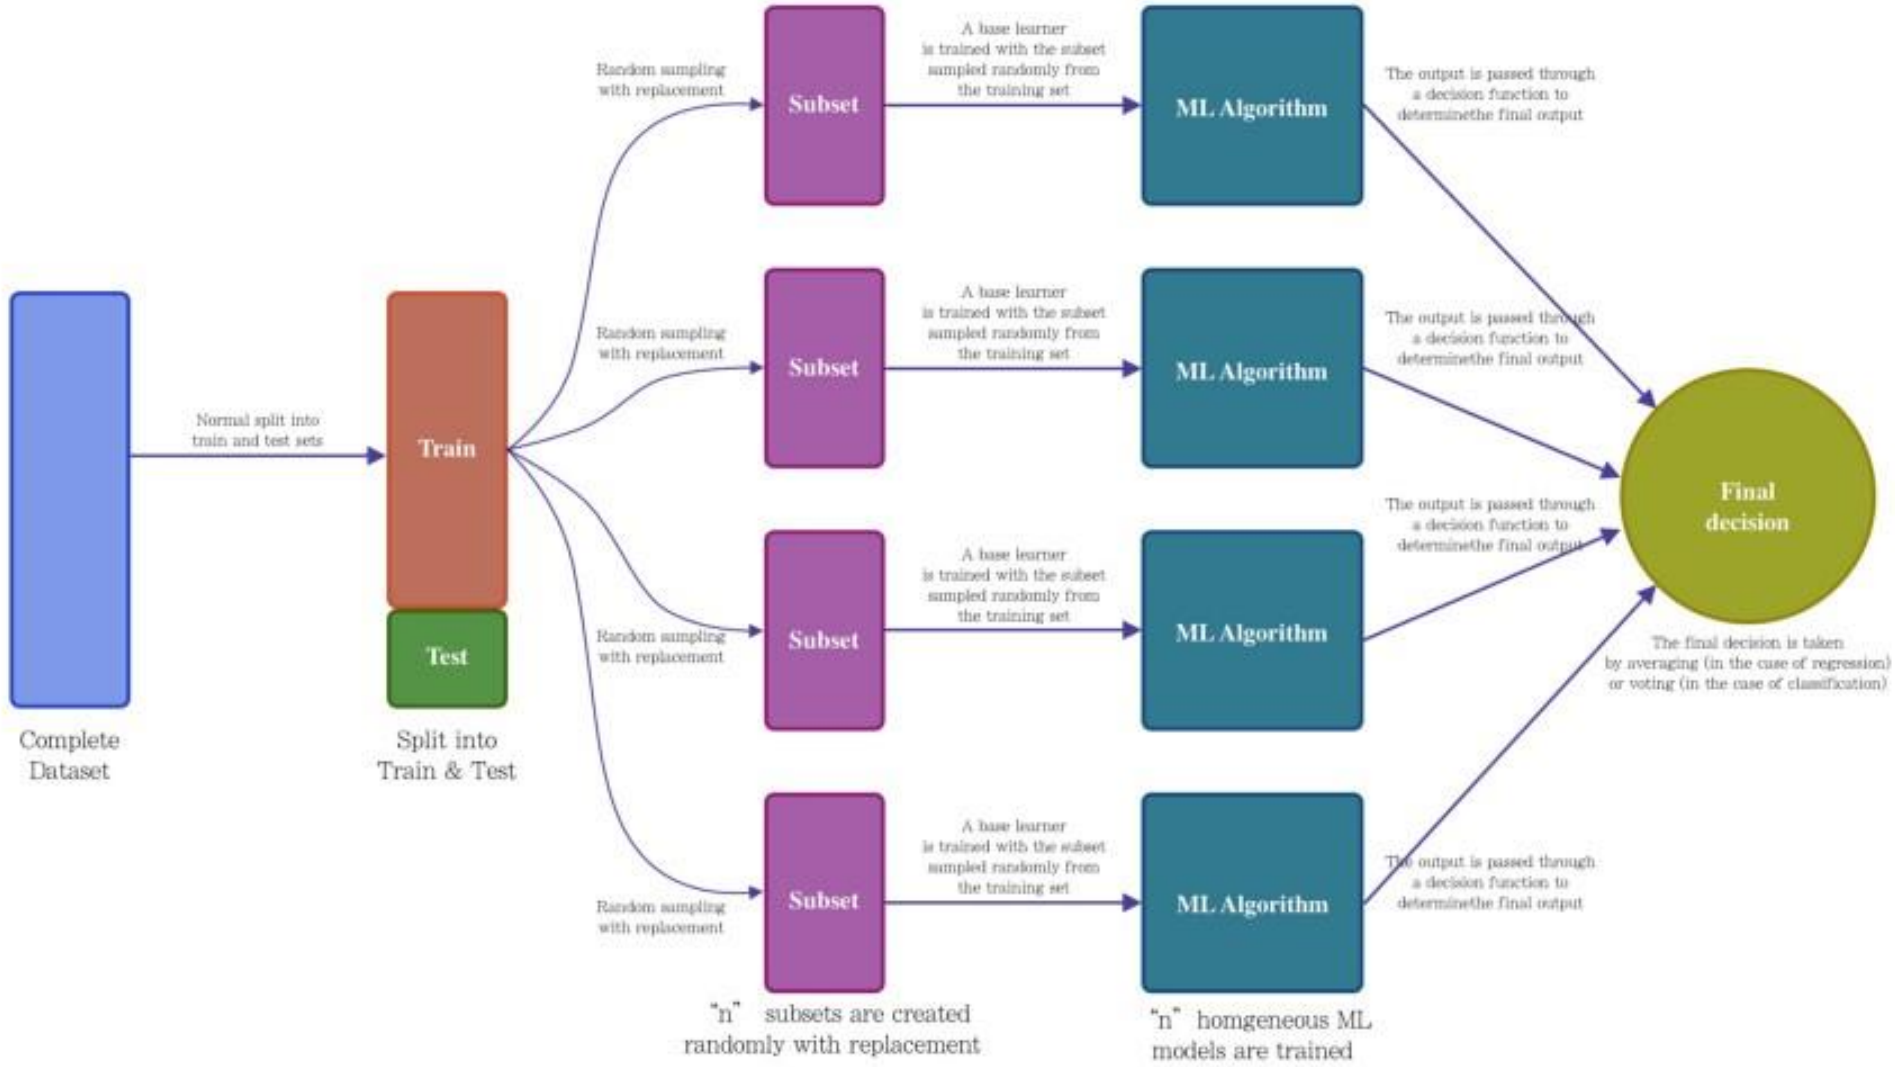
Complete Dataset
Normal split into train and test sets
Train
Test
Split into Train & Test
Random sampling with replacement
Subset
A base learner is trained with the subset sampled randomly from the training set
ML Algorithm
The output is passed through a decision function to determinethe final output
"n" subsets are created randomly with replacement
"n" homgeneous ML models are trained
Final decision
The final decision is taken by averaging (in the case of regression) or voting (in the case of classification)

Random Forest

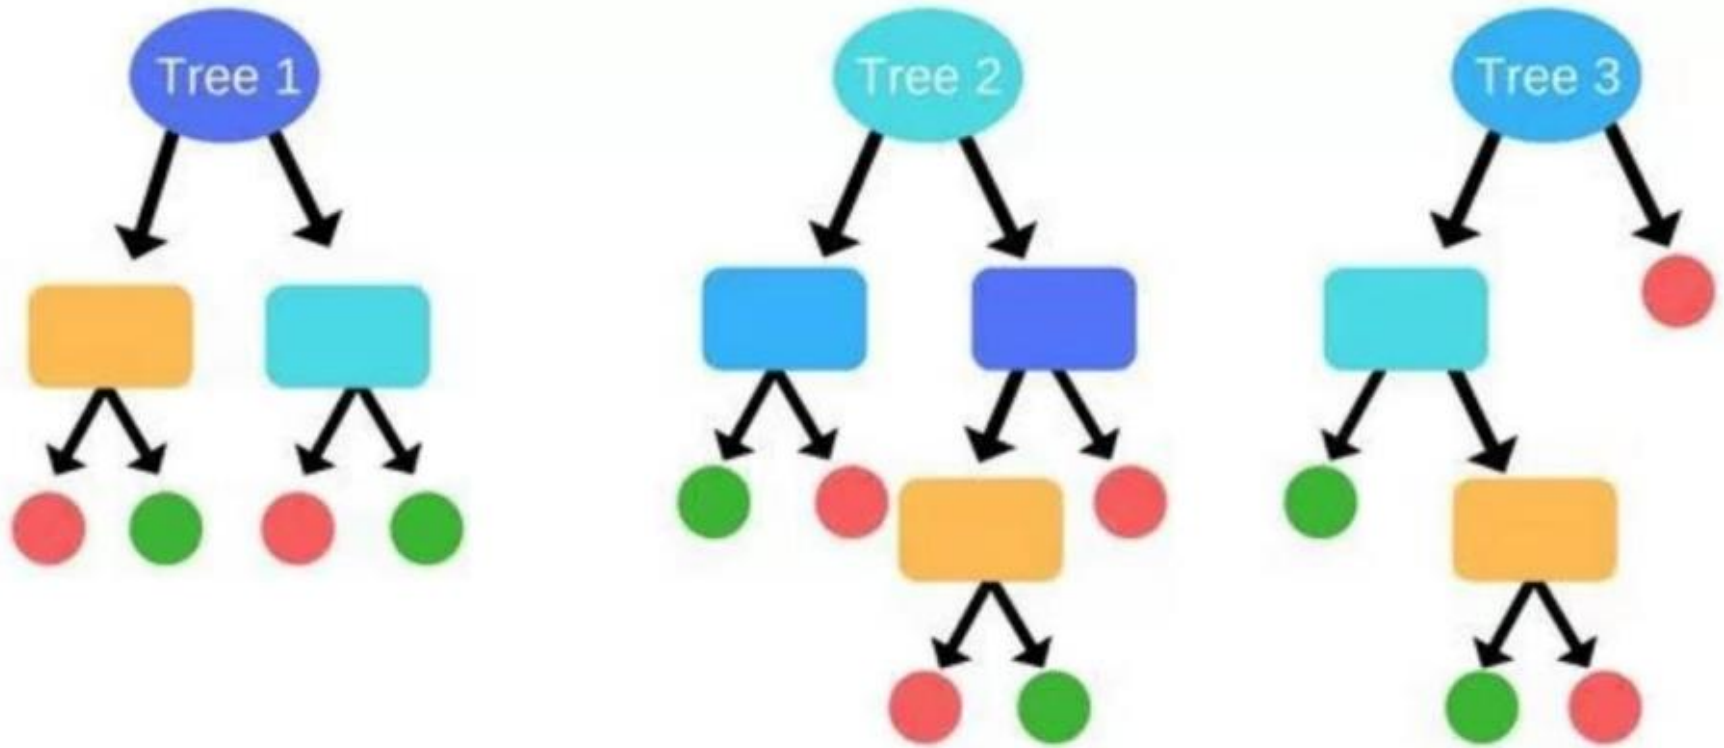

Boosting

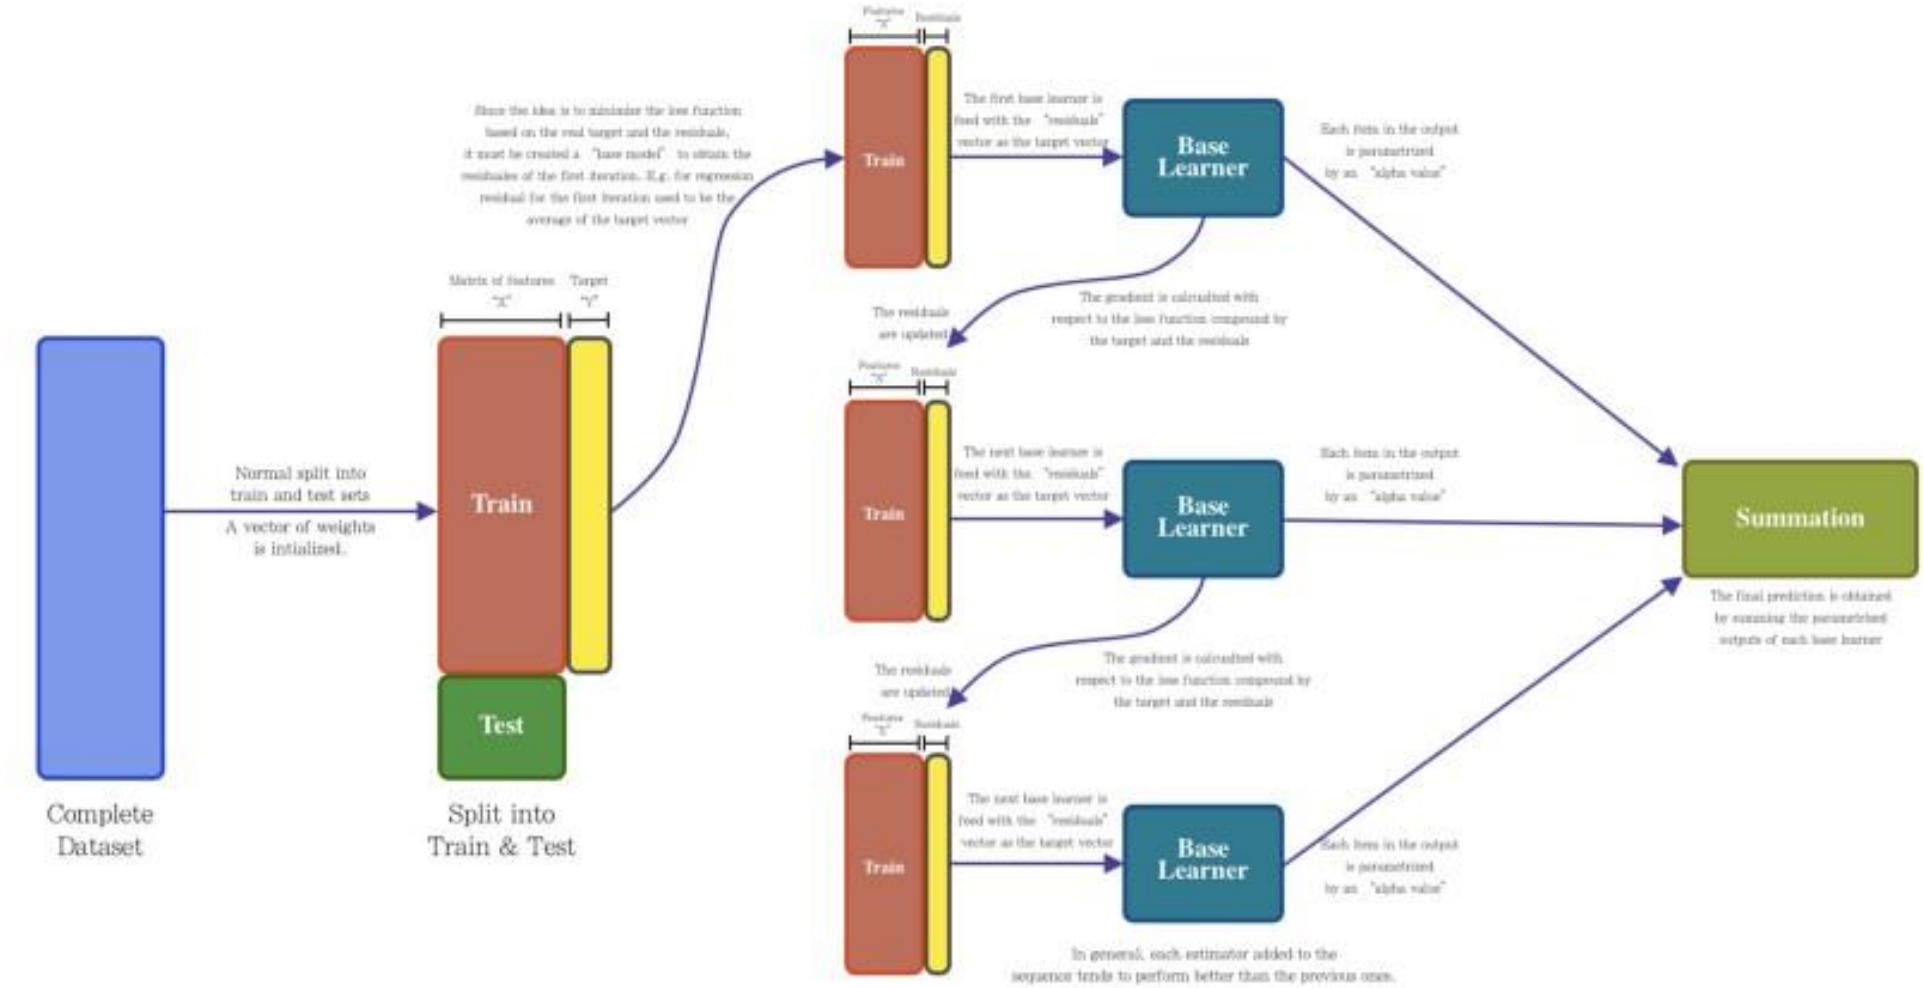
Complete Dataset
Normal split into train and test sets
A vector of weights is intialized.
Train
Test
Split into Train & Test
Train
Base Learner
Train
Base Learner
Train
Base Learner
Summation

Random Forest vs Gradient Boosting

Random Forest	Gradient Boosting
Easier to tune Harder to overfit Show very low variance Easier to parallelize	Better accuracy with less trees if the data is noisy exhibits higher variance

Ensembling summary